

Project : Customer Churn Dataset

1. Data Manipulation:

- Extract the 5th column and store it in 'customer_5'

Code:-

```
import pandas as pd

# Load the dataset (example: 'data.csv')
data = pd.read_csv('/content/customer_churn (1).csv')

# Extract the 5th column and store it in 'customer_5'
customer_5 = data.iloc[:, 4] # Indexing starts from 0, so the 5th column is index 4

# Print or verify the extracted column
print(customer_5)
```

Ans:-

```
0      No
1      No
2      No
3      No
4      No
...
7038   Yes
7039   Yes
7040   Yes
7041   No
7042   No
Name: Dependents, Length: 7043, dtype: object
```

- Extract the 15th column and store it in 'customer_15'

Code:-

```
import pandas as pd

# Load the dataset (example: 'data.csv')
data = pd.read_csv('/content/customer_churn (1).csv')

# Extract the 15th column and store it in 'customer_15'
customer_15 = data.iloc[:, 14] # Indexing starts from 0, so the 15th column is index 14

# Print or verify the extracted column
print(customer_15)
```

Ans :-

```
0      No
1      No
2      No
3      No
4      No
...
7038   Yes
7039   Yes
7040   No
7041   No
7042   Yes
Name: StreamingMovies, Length: 7043, dtype: object
```

- Extract all the male senior citizens whose payment method is electronic check and store the result in 'senior_male_electronic'

Code:-

```
import pandas as pd

# Load the dataset (example: 'data.csv')
data = pd.read_csv('/content/customer_churn (1).csv')

# Filter male senior citizens with payment method "Electronic check"
senior_male_electronic = data[(data['gender'] == 'Male') &
                              (data['SeniorCitizen'] == 1) &
                              (data['PaymentMethod'] == 'Electronic check')]

# Print or verify the filtered result
print(senior_male_electronic)
```

Ans :-

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	\
20	8779-QRDMV	Male	1	No	No	1	
55	1658-BYGOY	Male	1	No	No	18	
57	5067-XJQFU	Male	1	Yes	Yes	66	
78	0191-ZHSKZ	Male	1	No	No	30	
91	2424-WVHPL	Male	1	No	No	1	
...	
6837	6229-LSCKB	Male	1	No	No	6	
6894	1400-MMYXY	Male	1	Yes	No	3	
6914	7142-HVGBG	Male	1	Yes	No	43	
6967	8739-WWKDU	Male	1	No	No	25	
7032	6894-LFHLV	Male	1	No	No	1	

	PhoneService	MultipleLines	InternetService	OnlineSecurity	...	\
20	No	No phone service	DSL	No	...	
55	Yes	Yes	Fiber optic	No	...	
57	Yes	Yes	Fiber optic	No	...	
78	Yes	No	DSL	Yes	...	
91	Yes	No	Fiber optic	No	...	
...	
6837	Yes	No	Fiber optic	No	...	
6894	Yes	Yes	Fiber optic	No	...	
6914	Yes	Yes	Fiber optic	No	...	
6967	Yes	Yes	Fiber optic	No	...	
7032	Yes	Yes	Fiber optic	No	...	

	DeviceProtection	TechSupport	StreamingTV	StreamingMovies	Contract	\
20	Yes	No	No	Yes	Month-to-month	
55	No	No	Yes	Yes	Month-to-month	
57	Yes	Yes	Yes	Yes	One year	
78	No	No	Yes	Yes	Month-to-month	
91	No	Yes	No	No	Month-to-month	
...	
6837	No	No	Yes	No	Month-to-month	
6894	Yes	No	Yes	Yes	Month-to-month	
6914	Yes	No	Yes	Yes	Month-to-month	
6967	No	No	Yes	No	Month-to-month	
7032	No	No	No	No	Month-to-month	

	PaperlessBilling	PaymentMethod	MonthlyCharges	TotalCharges	Churn
20	Yes	Electronic check	39.65	39.65	Yes
55	Yes	Electronic check	95.45	1752.55	Yes
57	Yes	Electronic check	108.45	7076.35	No
78	Yes	Electronic check	74.75	2111.3	No
91	No	Electronic check	74.70	74.7	No
...
6837	Yes	Electronic check	79.70	497.6	No
6894	Yes	Electronic check	105.90	334.65	Yes
6914	Yes	Electronic check	103.00	4414.3	Yes
6967	Yes	Electronic check	89.50	2196.15	Yes
7032	Yes	Electronic check	75.75	75.75	Yes

[298 rows x 21 columns]

- Extract all those customers whose tenure is greater than 70 months or their monthly charges is more than \$100 and store the result in 'customer_total_tenure'

Code:-

```
import pandas as pd

# Sample data (replace this with your actual data loading step)
data = {
    'CustomerID': [1, 2, 3, 4, 5],
    'Tenure': [50, 80, 60, 75, 40],
    'MonthlyCharges': [90, 110, 70, 120, 100]
}
df = pd.DataFrame(data)

# Extract customers meeting the conditions
customer_total_tenure = df[(df['Tenure'] > 70) | (df['MonthlyCharges'] > 100)]

# Display the result
print(customer_total_tenure)
```

Ans :-

	CustomerID	Tenure	MonthlyCharges
1	2	80	110
3	4	75	120

- Extract all the customers whose contract is of two years, payment method is mailed check and the value of churn is 'Yes' and store the result in 'two_mail_yes'

Code :-

```
import pandas as pd

# Sample data (replace this with your actual data loading step)
data = {
    'CustomerID': [1, 2, 3, 4, 5],
    'Contract': ['Two years', 'Month-to-month', 'Two years', 'Two years', 'One year'],
    'PaymentMethod': ['Mailed check', 'Electronic check', 'Mailed check', 'Mailed check', 'Bank transfer'],
    'Churn': ['No', 'Yes', 'Yes', 'No', 'Yes']
}
df = pd.DataFrame(data)

# Extract customers meeting the conditions
two_mail_yes = df[
    (df['Contract'] == 'Two years') &
    (df['PaymentMethod'] == 'Mailed check') &
    (df['Churn'] == 'Yes')
]

# Display the result
print(two_mail_yes)
```

```
CustomerID  Contract  PaymentMethod  Churn
2           3   Two years   Mailed check   Yes
```

Ans :- CustomerID Contract PaymentMethod Churn

 2 3 Two Year Mailed Check Yes

- Extract 333 random records from the customer_churndataframe and store the result in 'customer_333'

Code:-

```
import pandas as pd

# Assuming customer_churndataframe is already loaded
# For demonstration, replace this with your actual DataFrame
customer_churndataframe = pd.DataFrame({
    'CustomerID': range(1, 1001), # Example data
    'Feature': ['SampleData'] * 1000
})

# Extract 333 random records
customer_333 = customer_churndataframe.sample(n=333, random_state=42) # random_state for reproducibility

# Display the result
print(customer_333)
```

Ans :-

	CustomerID	Feature
521	522	SampleData
737	738	SampleData
740	741	SampleData
660	661	SampleData
411	412	SampleData
..
711	712	SampleData
133	134	SampleData
703	704	SampleData
311	312	SampleData
722	723	SampleData

[333 rows x 2 columns]

- Get the count of different levels from the 'Churn' column

Code:-

```
import pandas as pd

# Sample data (replace this with your actual data)
data = {
    'CustomerID': [1, 2, 3, 4, 5],
    'Churn': ['Yes', 'No', 'No', 'Yes', 'No']
}
df = pd.DataFrame(data)

# Get the count of different levels in the 'Churn' column
churn_counts = df['Churn'].value_counts()

# Display the result
print(churn_counts)
```

Ans :-

```
Churn
No     3
Yes    2
Name: count, dtype: int64
```

2. Data Visualization:

- Build a bar-plot for the 'InternetService' column:

Code :-

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

# Sample data (replace this with your actual data)
data = {
    'CustomerID': [1, 2, 3, 4, 5, 6],
    'InternetService': ['DSL', 'Fiber optic', 'No', 'DSL', 'Fiber optic', 'No']
}
df = pd.DataFrame(data)

# Count the occurrences of each category in the 'InternetService' column
internet_service_counts = df['InternetService'].value_counts()

# Create a bar plot
plt.figure(figsize=(8, 5))
sns.barplot(x=internet_service_counts.index, y=internet_service_counts.values, palette='viridis')


# Customize the plot
plt.title('Bar Plot of Internet Service Types', fontsize=16)
plt.xlabel('Internet Service', fontsize=12)
plt.ylabel('Count', fontsize=12)
plt.xticks(fontsize=10)
plt.yticks(fontsize=10)
plt.show()
```

Ans :-

<ipython-input-7-3128da88bacb>:17: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x=internet_service_counts.index, y=internet_service_counts.values, palette='viridis')
```



a. Set x-axis label to 'Categories of Internet Service'

Code:-

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

# Sample data (replace this with your actual data)
data = {
    'CustomerID': [1, 2, 3, 4, 5, 6],
    'InternetService': ['DSL', 'Fiber optic', 'No', 'DSL', 'Fiber optic', 'No']
}
df = pd.DataFrame(data)

# Count the occurrences of each category in the 'InternetService' column
internet_service_counts = df['InternetService'].value_counts()

# Create a bar plot
plt.figure(figsize=(8, 5))
sns.barplot(x=internet_service_counts.index, y=internet_service_counts.values, palette='viridis')

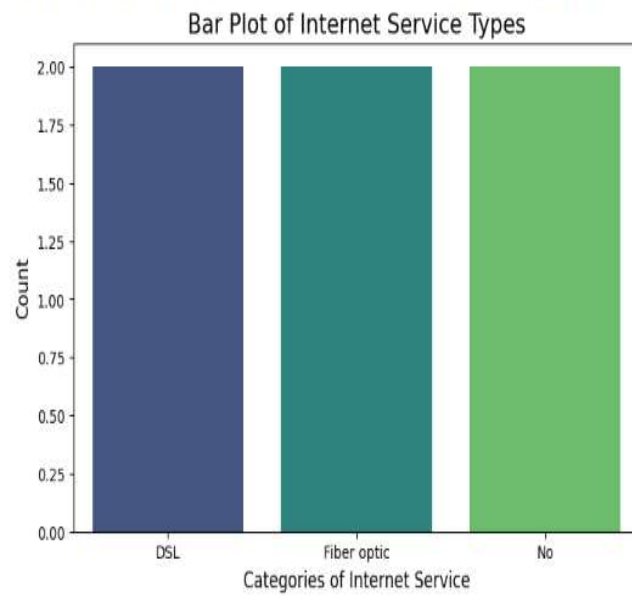
# Customize the plot
plt.title('Bar Plot of Internet Service Types', fontsize=16)
plt.xlabel('Categories of Internet Service', fontsize=12) # Set the x-axis label
plt.ylabel('Count', fontsize=12)
plt.xticks(fontsize=10)
plt.yticks(fontsize=10)
plt.show()
```

Ans ;-

```
<ipython-input-8-2fd4b0fc7948>:17: FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x=internet_service_counts.index, y=internet_service_counts.values, palette='viridis')
```



b. Set y-axis label to 'Count of Categories'

Code:-

```
import matplotlib.pyplot as plt

# Example data
categories = ['A', 'B', 'C', 'D']
counts = [10, 15, 20, 25]


# Creating the bar chart
plt.bar(categories, counts)

# Setting the y-axis label
plt.ylabel('Count of Categories')

#
```

Ans:-

Text(0, 0.5, 'Count of Categories')



c. Set the title of plot to be 'Distribution of Internet Service'

Code:-

```
import matplotlib.pyplot as plt
import numpy as np

# Generate some example data
np.random.seed(42)
data = np.random.normal(loc=5, scale=2, size=1000) # example normal distribution

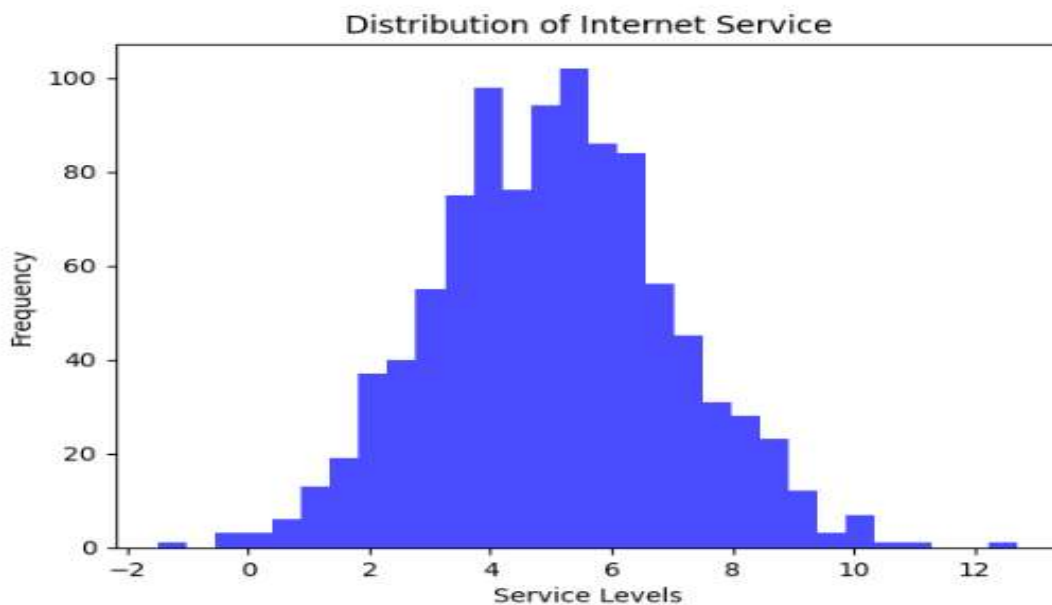
# Create a histogram
plt.hist(data, bins=30, color='blue', alpha=0.7)

# Set the title
plt.title('Distribution of Internet Service')

# Labeling the axes (optional)
plt.xlabel('Service Levels')
plt.ylabel('Frequency')

# Display the plot
plt.show()
```

Ans:-



d. Set the color of the bars to be 'orange'

Code:-

```
import matplotlib.pyplot as plt
import numpy as np

# Generate some example data
np.random.seed(42)
data = np.random.normal(loc=5, scale=2, size=1000) # example normal distribution


# Create a histogram with orange bars
plt.hist(data, bins=30, color='orange', alpha=0.7)

# Set the title
plt.title('Distribution of Internet Service')

# Labeling the axes (optional)
plt.xlabel('Service Levels')
plt.ylabel('Frequency')

# Display the plot
plt.show()
```

Ans :-



- Build a histogram for the 'tenure' column:

a. Set the number of bins to be 30

Code:-

```
import pandas as pd
import matplotlib.pyplot as plt

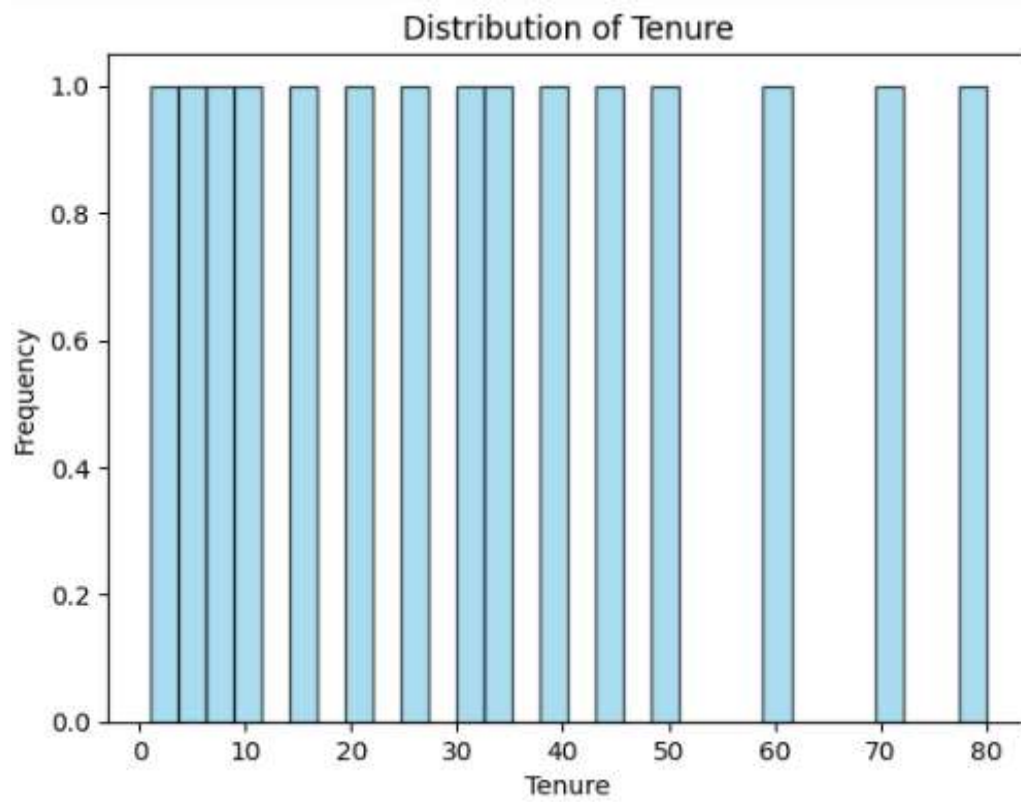
# Example DataFrame (replace with your actual dataset)
data = {'tenure': [1, 5, 8, 10, 15, 20, 25, 30, 35, 40, 45, 50, 60, 70, 80]}
df = pd.DataFrame(data)

# Create a histogram for the 'tenure' column with 30 bins
plt.hist(df['tenure'], bins=30, color='skyblue', edgecolor='black', alpha=0.7)

# Set title and labels
plt.title('Distribution of Tenure')
plt.xlabel('Tenure')
plt.ylabel('Frequency')

# Display the plot
plt.show()
```

Ans:-



b. Set the color of the bins to be 'green'

Code:--

```
import pandas as pd
import matplotlib.pyplot as plt

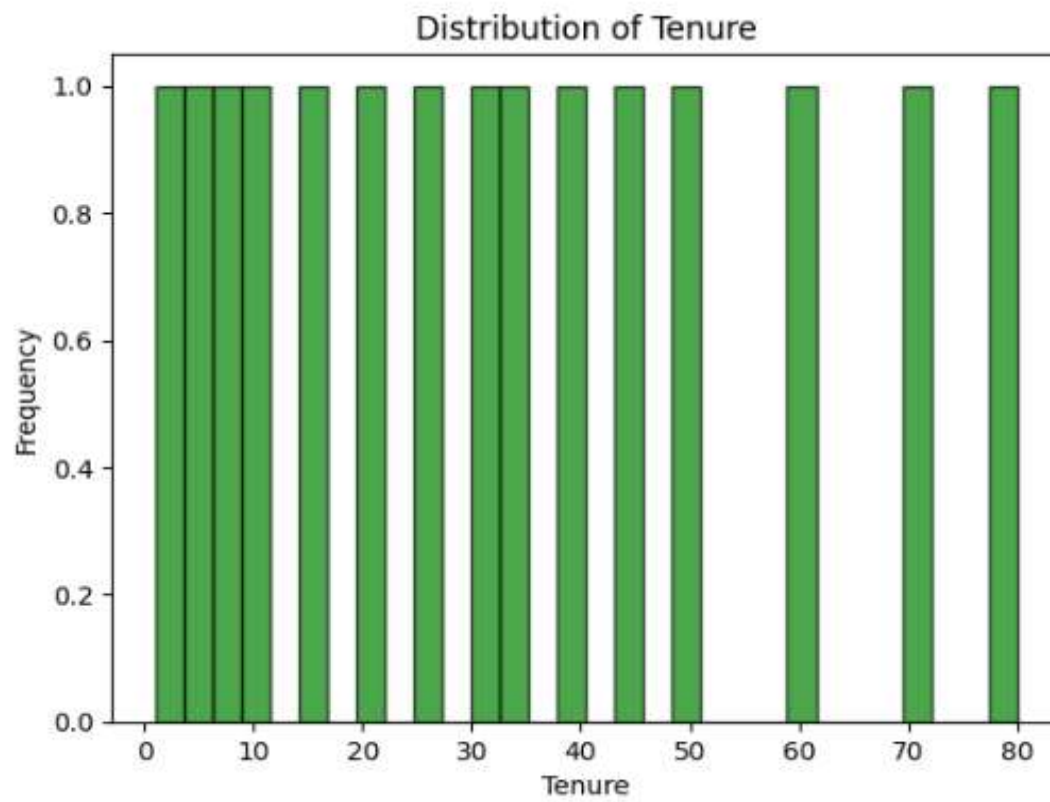
# Example DataFrame (replace with your actual dataset)
data = {'tenure': [1, 5, 8, 10, 15, 20, 25, 30, 35, 40, 45, 50, 60, 70, 80]}
df = pd.DataFrame(data)

# Create a histogram for the 'tenure' column with 30 bins and green color
plt.hist(df['tenure'], bins=30, color='green', edgecolor='black', alpha=0.7)

# Set title and labels
plt.title('Distribution of Tenure')
plt.xlabel('Tenure')
plt.ylabel('Frequency')

# Display the plot
plt.show()
```

Ans:-



c. Assign the title 'Distribution of tenure'

Code:-

```
import pandas as pd
import matplotlib.pyplot as plt

# Example DataFrame (replace with your actual dataset)
data = {'tenure': [1, 5, 8, 10, 15, 20, 25, 30, 35, 40, 45, 50, 60, 70, 80]}
df = pd.DataFrame(data)


# Create a histogram for the 'tenure' column with 30 bins and green color
plt.hist(df['tenure'], bins=30, color='green', edgecolor='black', alpha=0.7)

# Set the title
plt.title('Distribution of tenure')

# Set labels for the axes
plt.xlabel('Tenure')
plt.ylabel('Frequency')

# Display the plot
plt.show()
```

Ans:-



- Build a scatter-plot between 'MonthlyCharges' and 'tenure'. Map 'MonthlyCharges' to the y-axis and 'tenure' to the 'x-axis':

a. Assign the points a color of 'brown'

Code:-

```
import matplotlib.pyplot as plt

# Sample data (replace with your dataset)
# Replace these lists with your actual data for MonthlyCharges and tenure
tenure = [1, 2, 3, 4, 5, 6] # Example data for tenure
monthly_charges = [20, 25, 30, 35, 40, 45] # Example data for MonthlyCharges

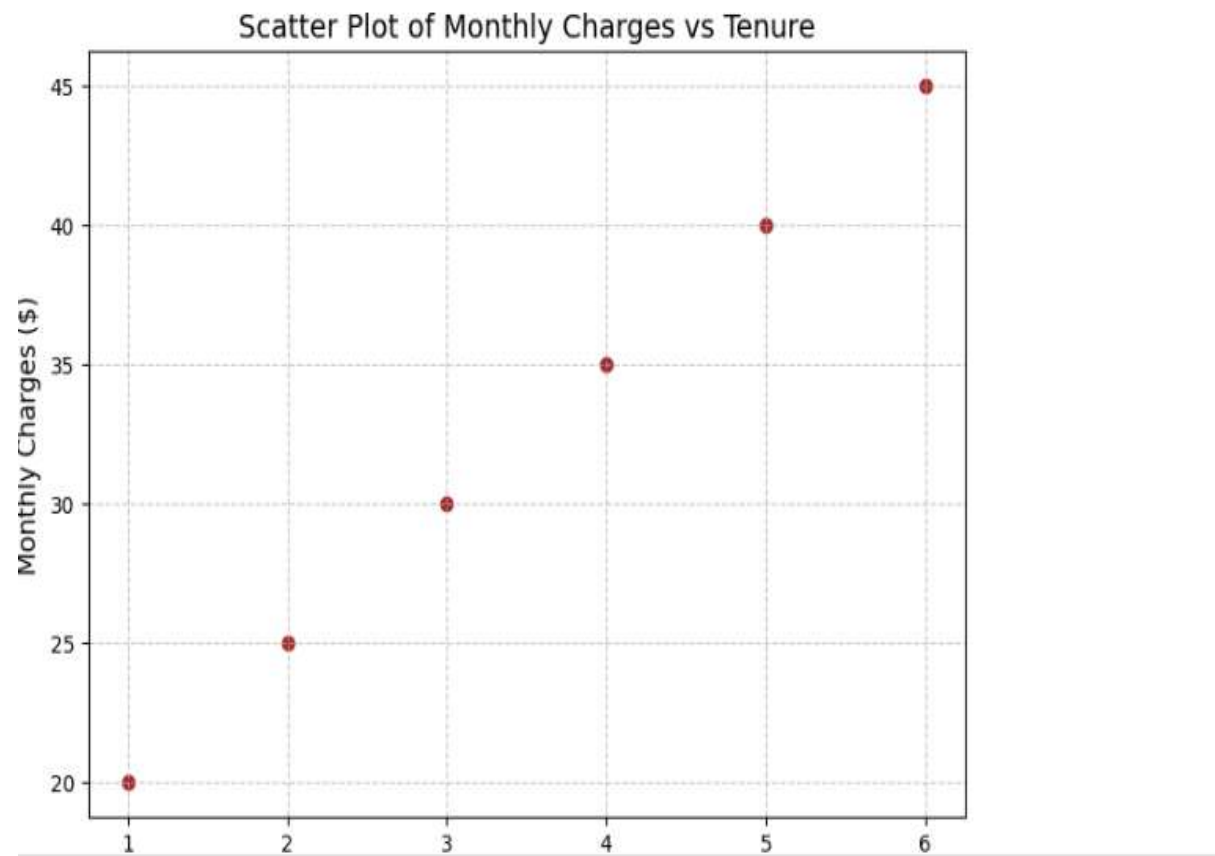
# Create the scatter plot
plt.figure(figsize=(8, 6))
plt.scatter(tenure, monthly_charges, color='brown')

# Label the axes
plt.xlabel('Tenure (months)', fontsize=12)
plt.ylabel('Monthly Charges ($)', fontsize=12)

# Title of the plot
plt.title('Scatter Plot of Monthly Charges vs Tenure', fontsize=14)

# Show the plot
plt.grid(True, linestyle='--', alpha=0.7)
plt.show()
```

Ans:-



b. Set the x-axis label to 'Tenure of customer'

Code:-

```
import matplotlib.pyplot as plt

# Sample data (replace with your dataset)
# Replace these lists with your actual data for MonthlyCharges and tenure
tenure = [1, 2, 3, 4, 5, 6] # Example data for tenure
monthly_charges = [20, 25, 30, 35, 40, 45] # Example data for MonthlyCharges

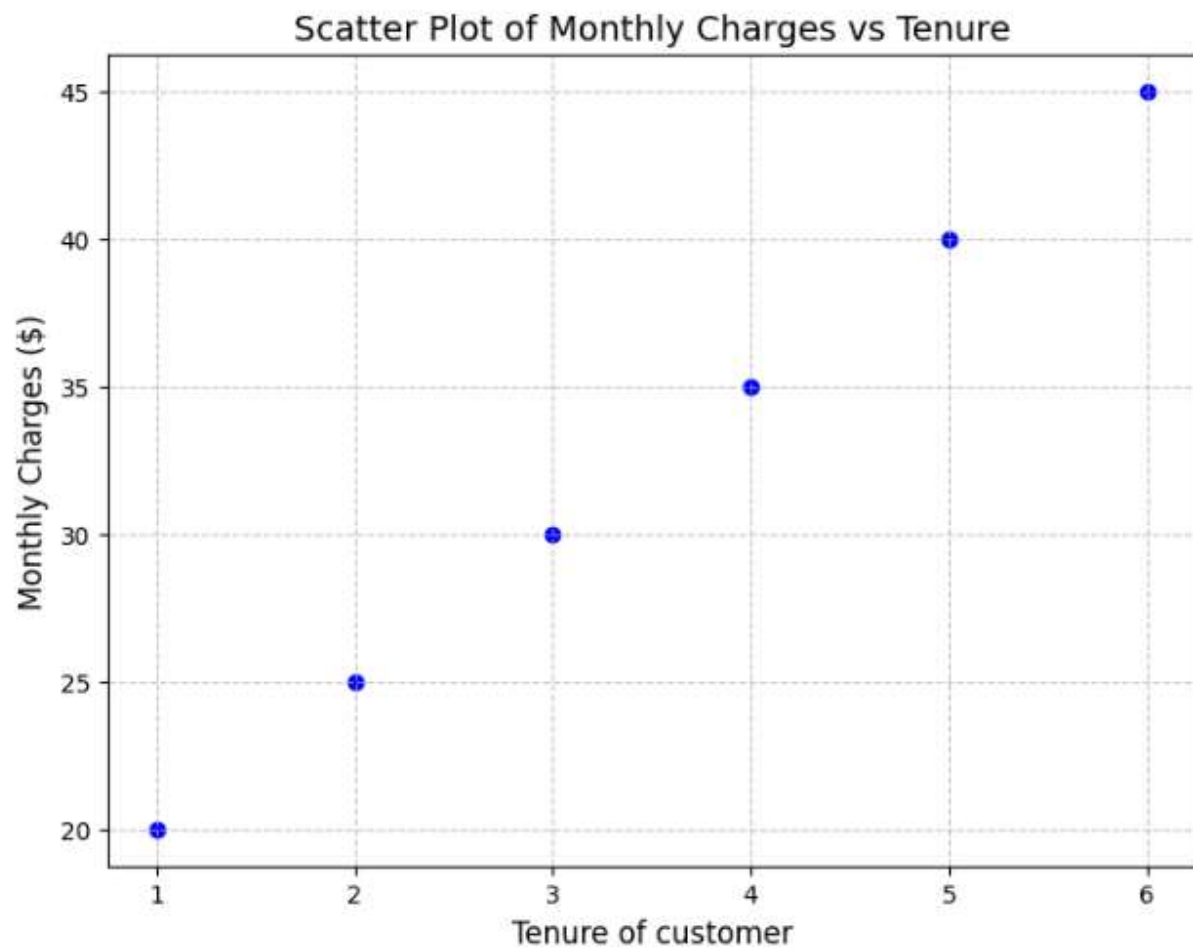
# Create the scatter plot
plt.figure(figsize=(8, 6))
plt.scatter(tenure, monthly_charges, color='blue') # Updated color to 'blue'

# Label the axes
plt.xlabel('Tenure of customer', fontsize=12)
plt.ylabel('Monthly Charges ($)', fontsize=12)

# Title of the plot
plt.title('Scatter Plot of Monthly Charges vs Tenure', fontsize=14)

# Show the plot
plt.grid(True, linestyle='--', alpha=0.7)
plt.show()
```

Ans:-



c. Set the y-axis label to 'Monthly Charges of customer'

Code:-

```
import matplotlib.pyplot as plt

# Sample data (replace with your dataset)
# Replace these lists with your actual data for MonthlyCharges and tenure
tenure = [1, 2, 3, 4, 5, 6] # Example data for tenure
monthly_charges = [20, 25, 30, 35, 40, 45] # Example data for MonthlyCharges


# Create the scatter plot
plt.figure(figsize=(8, 6))
plt.scatter(tenure, monthly_charges, color='red') # You can change the color if needed

# Label the axes
plt.xlabel('Tenure of customer', fontsize=12)
plt.ylabel('Monthly Charges of customer', fontsize=12) # Updated y-axis label

# Title of the plot
plt.title('Scatter Plot of Monthly Charges vs Tenure', fontsize=14)

# Show the plot
plt.grid(True, linestyle='--', alpha=0.7)
plt.show()
```

Ans:-



d. Set the title to 'Tenure vs Monthly Charges'

Code:-

```
import matplotlib.pyplot as plt

# Sample data (replace with your dataset)
# Replace these lists with your actual data for MonthlyCharges and tenure
tenure = [1, 2, 3, 4, 5, 6] # Example data for tenure
monthly_charges = [20, 25, 30, 35, 40, 45] # Example data for MonthlyCharges

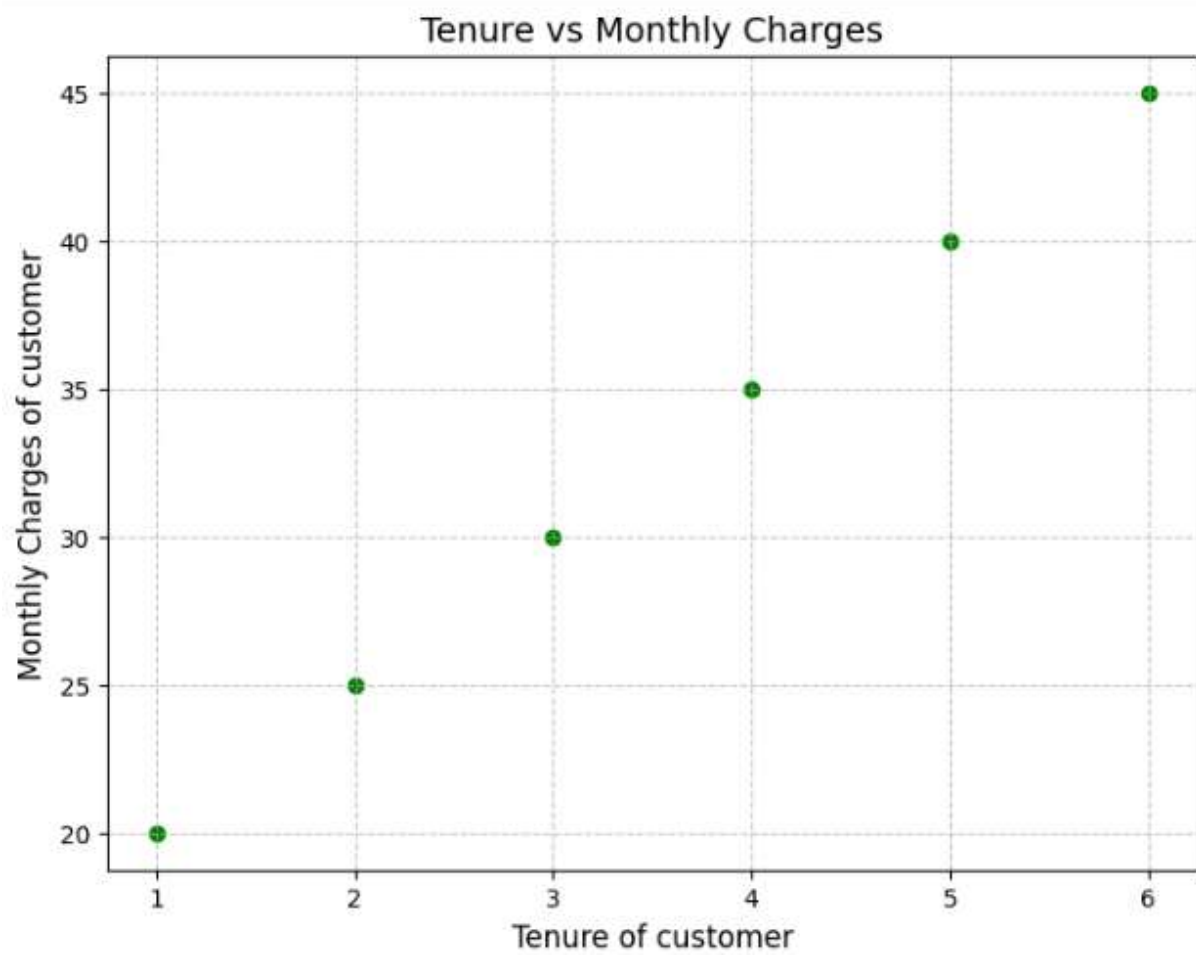
# Create the scatter plot
plt.figure(figsize=(8, 6))
plt.scatter(tenure, monthly_charges, color='green') # Updated color to 'green'

# Label the axes
plt.xlabel('Tenure of customer', fontsize=12)
plt.ylabel('Monthly Charges of customer', fontsize=12)

# Title of the plot
plt.title('Tenure vs Monthly Charges', fontsize=14)

# Show the plot
plt.grid(True, linestyle='--', alpha=0.7)
plt.show()
```

Ans:-



e. Build a box-plot between 'tenure' & 'Contract'. Map 'tenure' on the y-axis &

Code:-

```
import seaborn as sns
import matplotlib.pyplot as plt
import pandas as pd

# Sample data (replace with your dataset)
data = {
    'tenure': [1, 12, 24, 36, 48, 60, 72], # Example data for tenure
    'Contract': ['Month-to-month', 'One year', 'Month-to-month', 'Two year', 'One year', 'Two year', 'Month-to-month'] # Example data for Contract
}

# Convert to a DataFrame
df = pd.DataFrame(data)

# Create the box plot
plt.figure(figsize=(8, 6))
sns.boxplot(x='Contract', y='tenure', data=df, palette='Set2') # 'Set2' gives a nice color palette

# Label the axes and title
plt.xlabel('Contract Type', fontsize=12)
plt.ylabel('Tenure of Customer (months)', fontsize=12)
plt.title('Box Plot of Tenure vs Contract Type', fontsize=14)

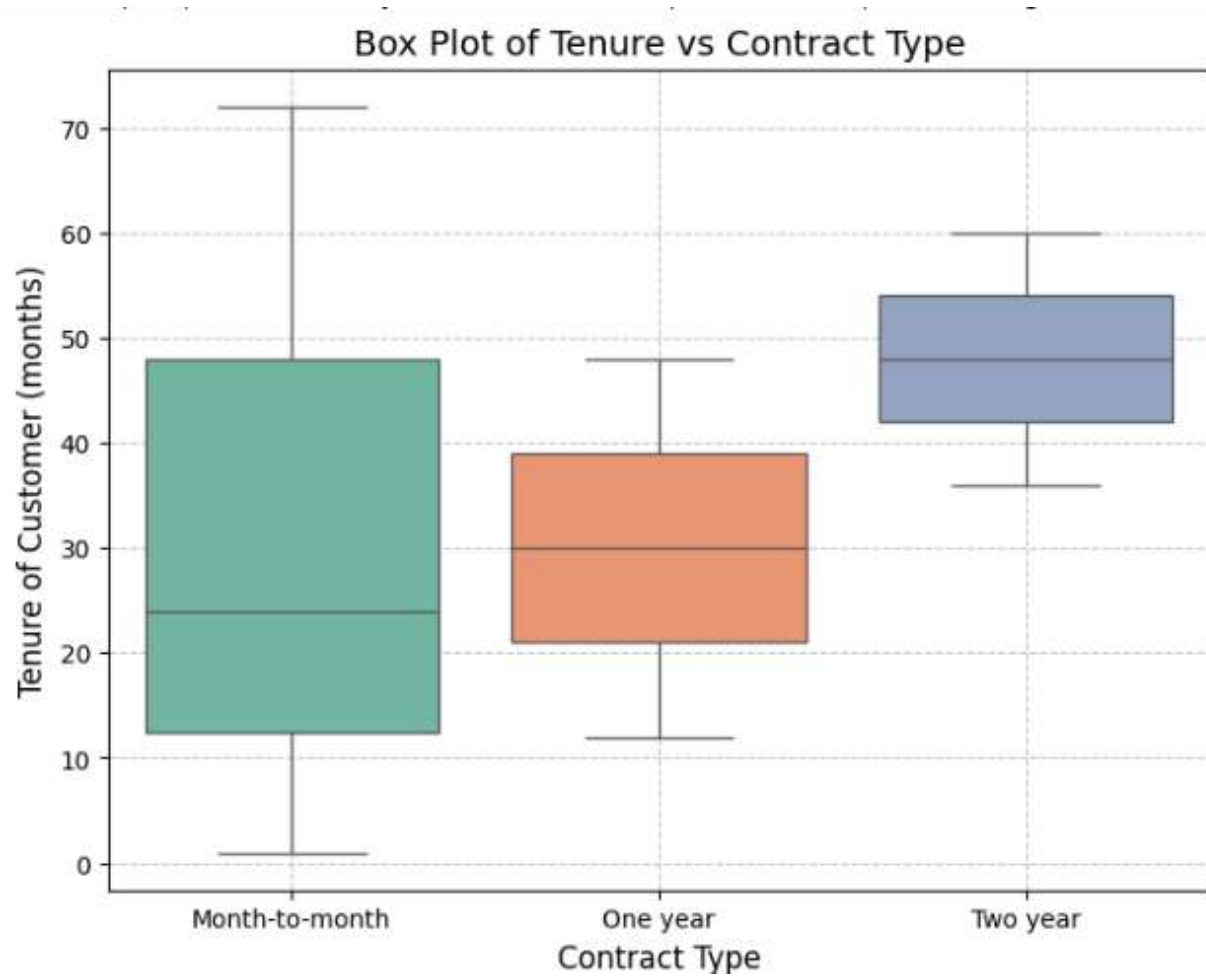
# Show the plot
plt.grid(True, linestyle='--', alpha=0.7)
plt.show()
```

Ans:-

```
<ipython-input-10-51c446351dc4>:16: FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(x='Contract', y='tenure', data=df, palette='Set2') # 'Set2' gives a nice color palette
```



3. Linear Regression:

- Build a simple linear model where dependent variable is 'MonthlyCharges' and independent variable is 'tenure':

a. Divide the dataset into train and test sets in 70:30 ratio

Code:-

```
# Import necessary libraries
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score

# Load the dataset (replace 'your_dataset.csv' with your actual dataset file)
data = pd.read_csv('/content/customer_churn (1).csv')

# Select the dependent and independent variables
X = data[['tenure']] # Independent variable
y = data['MonthlyCharges'] # Dependent variable

# Split the dataset into training and testing sets (70:30 ratio)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

# Initialize and train the linear regression model
model = LinearRegression()
model.fit(X_train, y_train)

# Predict on the test set
y_pred = model.predict(X_test)

# Evaluate the model
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

# Display results
print(f"Mean Squared Error (MSE): {mse}")
print(f"R-squared (R2): {r2}")
```

```
# Display model coefficients
print(f"Intercept: {model.intercept_}")
print(f"Coefficient for 'tenure': {model.coef_[0]}")
```

Ans:-

Mean Squared Error (MSE): 845.6091871095867
R-squared (R2): 0.05856035027031625
Intercept: 54.79837462739951
Coefficient for 'tenure': 0.3082154776200297

b. Build the model on train set and predict the values on test set

Code:-

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score

# Load dataset
data = pd.read_csv('/content/customer_churn (1).csv')
X, y = data[['tenure']], data['MonthlyCharges']

# Split data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

# Train and predict
model = LinearRegression().fit(X_train, y_train)
y_pred = model.predict(X_test)

# Evaluate
print(f"MSE: {mean_squared_error(y_test, y_pred)}")
print(f"R²: {r2_score(y_test, y_pred)}")
print(f"Intercept: {model.intercept_}, Coefficient: {model.coef_[0]}")
```

Ans:-

MSE: 845.6091871095867
R²: 0.05856035027031625
Intercept: 54.79837462739951, Coefficient: 0.3082154776200297

c. After predicting the values, find the root mean square error

Code:-

```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error

X, y = data[['tenure']], data['MonthlyCharges']

# Split data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

# Train and predict
model = LinearRegression().fit(X_train, y_train)
y_pred = model.predict(X_test)

# Calculate RMSE manually
rmse = np.sqrt(mean_squared_error(y_test, y_pred))
print(f"Root Mean Square Error (RMSE): {rmse}")
```

Asn:-

Root Mean Square Error (RMSE): 29.07936015646814

d. Find out the error in prediction & store the result in 'error'

Code:-

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression

X, y = data[['tenure']], data['MonthlyCharges']

# Split data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

# Train and predict
model = LinearRegression().fit(X_train, y_train)
y_pred = model.predict(X_test)

# Calculate error and store in 'error'
error = y_test - y_pred

# Display errors
print("Prediction Errors:")
print(error)
```

Ans:-

```
Prediction Errors:
185      -30.306590
2715     -42.185209
3825     -51.475579
1807      21.243410
132      -24.898812
...
5522      15.043410
6377     -10.456590
5500     -12.791303
2392      29.737316
6705     -10.479115
Name: MonthlyCharges, Length: 2113, dtype: float64
```

e. Find the root mean square error

Code:-

```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error

X, y = data[['tenure']], data['MonthlyCharges']

# Split data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

# Train and predict
model = LinearRegression().fit(X_train, y_train)
y_pred = model.predict(X_test)

# Calculate RMSE
rmse = np.sqrt(mean_squared_error(y_test, y_pred))
print(f"Root Mean Square Error (RMSE): {rmse}")
```

Ans:-

Root Mean Square Error (RMSE): 29.07936015646814

4. Logistic Regression:

- Build a simple logistic regression model where dependent variable is 'Churn' and independent variable is 'MonthlyCharges':

a. Divide the dataset in 65:35 ratio

Code:-

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report

# Select the dependent and independent variables
X = data[['MonthlyCharges']] # Independent variable
y = data['Churn'] # Dependent variable (ensure this is binary: 0 or 1)

# Split dataset into training and test sets (65:35 ratio)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.35, random_state=42)

# Initialize and train the logistic regression model
model = LogisticRegression()
model.fit(X_train, y_train)

# Make predictions on the test set
y_pred = model.predict(X_test)

# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy}")
print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))
print("Classification Report:")
print(classification_report(y_test, y_pred))
```


Ans:-

```
Accuracy: 0.7287104622871047
Confusion Matrix:
[[1797  0]
 [ 669  0]]
Classification Report:
      precision    recall  f1-score   support

     No         0.73      1.00      0.84       1797
     Yes         0.00      0.00      0.00        669

 accuracy          0.73      2466
 macro avg         0.36      0.50      0.42      2466
 weighted avg      0.53      0.73      0.61      2466

/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples
_warn_prf(average, modifier, f'{metric.capitalize()} is', len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples
_warn_prf(average, modifier, f'{metric.capitalize()} is', len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples
_warn_prf(average, modifier, f'{metric.capitalize()} is', len(result))
```

b. Build the model on train set and predict the values on test set

code:-

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score

# Select the dependent and independent variables
X = data[['MonthlyCharges']] # Independent variable
y = data['Churn'] # Dependent variable (binary: 0 or 1)

# Split dataset into training and test sets (65:35 ratio)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.35, random_state=42)

# Initialize and train the logistic regression model on the training set
model = LogisticRegression()
model.fit(X_train, y_train)

# Predict values on the test set
y_pred = model.predict(X_test)

# Display predictions
print("Predicted values on the test set:")
print(y_pred)

# Evaluate the model (optional)
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy}")
```

Ans:-

```
· Predicted values on the test set:
  ['No' 'No' 'No' ... 'No' 'No' 'No']
  Accuracy: 0.7287104622871047
```


c. Build the confusion matrix and get the accuracy score

Code:-

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix, accuracy_score

# Select the dependent and independent variables
X = data[['MonthlyCharges']] # Independent variable
y = data['Churn'] # Dependent variable (binary: 0 or 1)

# Split dataset into training and test sets (65:35 ratio)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.35, random_state=42)

# Initialize and train the logistic regression model
model = LogisticRegression()
model.fit(X_train, y_train)

# Predict values on the test set
y_pred = model.predict(X_test)

# Build the confusion matrix
conf_matrix = confusion_matrix(y_test, y_pred)
print("Confusion Matrix:")
print(conf_matrix)

# Calculate and display the accuracy score
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy Score: {accuracy}")
```

Ans:-

```
Confusion Matrix:
[[1797   0]
 [ 669   0]]
Accuracy Score: 0.7287104622871047
```

d. Build a multiple logistic regression model where dependent variable is 'Churn' and independent variables are 'tenure' and 'MonthlyCharges'

Code:-

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix, accuracy_score

# Select the dependent and independent variables
X = data[['tenure', 'MonthlyCharges']] # Independent variables
y = data['Churn'] # Dependent variable (binary: 0 or 1)

# Split dataset into training and test sets (65:35 ratio)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.35, random_state=42)

# Initialize and train the logistic regression model
model = LogisticRegression()
model.fit(X_train, y_train)

# Predict values on the test set
y_pred = model.predict(X_test)

# Build the confusion matrix
conf_matrix = confusion_matrix(y_test, y_pred)
print("Confusion Matrix:")
print(conf_matrix)

# Calculate and display the accuracy score
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy Score: {accuracy}")
```

Ans:-

```
Confusion Matrix:
[[1650  147]
 [ 359  310]]
Accuracy Score: 0.7948094079480941
```

e. Divide the dataset in 80:20 ratio

Code:-

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix, accuracy_score

# Define features and target
X = data[['tenure', 'MonthlyCharges']]
y = data['Churn']

# Split data (80:20 ratio)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Train logistic regression model
model = LogisticRegression()
model.fit(X_train, y_train)

# Predict and evaluate
y_pred = model.predict(X_test)
print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))
print(f"Accuracy Score: {accuracy_score(y_test, y_pred)}")
```

```
Confusion Matrix:
[[944  92]
 [193 180]]
Accuracy Score: 0.7977288857345636
```

Ans:-

Confusion Matrix:- [[944 92]

[193 180]]

Accuracy Score: 0.7977288857345636

5. Decision Tree:

- Build a decision tree model where dependent variable is 'Churn' and independent variable is 'tenure'.

a. Divide the dataset in 80:20 ratio

Code:-

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier, plot_tree
import matplotlib.pyplot as plt

# Load your dataset
df = pd.DataFrame({'tenure': [1, 2, 3, 4, 5], 'Churn': ['Yes', 'No', 'Yes', 'No', 'Yes']})

# Prepare data
X = df[['tenure']]
y = df['Churn'].map({'Yes': 1, 'No': 0}) # Convert 'Churn' to numeric

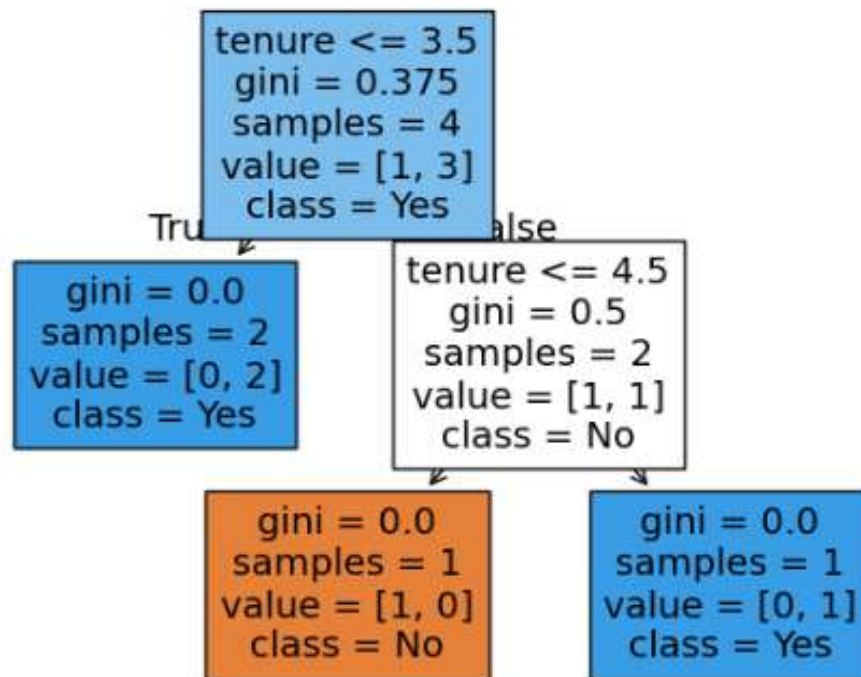
# Split data (80:20)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Train decision tree
clf = DecisionTreeClassifier(random_state=42)
clf.fit(X_train, y_train)

# Visualize tree
plot_tree(clf, feature_names=['tenure'], class_names=['No', 'Yes'], filled=True)
plt.show()

# Print accuracy
print(f"Accuracy: {clf.score(X_test, y_test):.2f}")
```

Ans:-



Accuracy: 0.00

b. Build the model on train set and predict the values on test set

Code:-

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score, classification_report

# Example dataset
data = {
    'tenure': [1, 2, 3, 4, 5, 6, 7, 8, 9, 10],
    'Churn': ['Yes', 'No', 'Yes', 'No', 'Yes', 'Yes', 'No', 'No', 'Yes', 'No']
}
df = pd.DataFrame(data)

# Map 'Churn' to numeric values
X = df[['tenure']] # Independent variable
y = df['Churn'].map({'Yes': 1, 'No': 0}) # Dependent variable

# Split data into training and test sets (80:20 split)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42, stratify=y)

# Train the Decision Tree model
clf = DecisionTreeClassifier(max_depth=3, random_state=42)
clf.fit(X_train, y_train)

# Predict values on the test set
y_pred = clf.predict(X_test)

# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy on Test Set: {accuracy:.2f}")
print("Classification Report:")
print(classification_report(y_test, y_pred, target_names=['No', 'Yes']))
```

Ans:-

Accuracy on Test Set: 0.50

Classification Report:

	precision	recall	f1-score	support
No	0.50	1.00	0.67	1
Yes	0.00	0.00	0.00	1
accuracy			0.50	2
macro avg	0.25	0.50	0.33	2
weighted avg	0.25	0.50	0.33	2

```
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples
  _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples
  _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples
  _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```


c. Build the confusion matrix and calculate the accuracy

Code:-

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import confusion_matrix, accuracy_score, classification_report

# Example dataset
data = {
    'tenure': [1, 2, 3, 4, 5, 6, 7, 8, 9, 10],
    'Churn': ['Yes', 'No', 'Yes', 'No', 'Yes', 'Yes', 'No', 'No', 'Yes', 'No']
}
df = pd.DataFrame(data)

# Map 'Churn' to numeric values
X = df[['tenure']] # Independent variable
y = df['Churn'].map({'Yes': 1, 'No': 0}) # Dependent variable

# Split data into training and test sets (80:20 split)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42, stratify=y)

# Train the Decision Tree model
clf = DecisionTreeClassifier(max_depth=3, random_state=42)
clf.fit(X_train, y_train)

# Predict values on the test set
y_pred = clf.predict(X_test)

# Build the confusion matrix
cm = confusion_matrix(y_test, y_pred)
print("Confusion Matrix:")
print(cm)

# Calculate accuracy
```

```
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy:.2f}")

# Optionally, print the classification report
print("Classification Report:")
print(classification_report(y_test, y_pred, target_names=['No', 'Yes']))
```


Ans:-

Confusion Matrix:

```
[[1 0]
 [1 0]]
```

Accuracy: 0.50

Classification Report:

	precision	recall	f1-score	support
No	0.50	1.00	0.67	1
Yes	0.00	0.00	0.00	1
accuracy			0.50	2
macro avg	0.25	0.50	0.33	2
weighted avg	0.25	0.50	0.33	2

```
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples
_warn_prf(average, modifier, f'{metric.capitalize()} is', len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples
_warn_prf(average, modifier, f'{metric.capitalize()} is', len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples
_warn_prf(average, modifier, f'{metric.capitalize()} is', len(result))
```

6. Random Forest:

- Build a Random Forest model where dependent variable is 'Churn' and independent variables are 'tenure' and 'MonthlyCharges':

a. Divide the dataset in 70:30 ratio

Code:-

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report

# Sample Data
df = pd.DataFrame({
    'tenure': [1, 2, 3, 4, 5, 6, 7, 8, 9, 10],
    'MonthlyCharges': [20, 30, 40, 50, 60, 70, 80, 90, 100, 110],
    'Churn': ['Yes', 'No', 'Yes', 'No', 'Yes', 'Yes', 'No', 'No', 'Yes', 'No']
})

# Prepare Data
X = df[['tenure', 'MonthlyCharges']]
y = df['Churn'].map({'Yes': 1, 'No': 0})
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, stratify=y, random_state=42)

# Train and Predict
clf = RandomForestClassifier(n_estimators=100, random_state=42)
clf.fit(X_train, y_train)
y_pred = clf.predict(X_test)

# Results
print(f"Accuracy: {accuracy_score(y_test, y_pred):.2f}")
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))
print("Classification Report:\n", classification_report(y_test, y_pred, target_names=['No', 'Yes']))
```

Ans:-

Accuracy: 0.33

Confusion Matrix:

[[1 1]

[1 0]]

Classification Report:

	precision	recall	f1-score	support
No	0.50	0.50	0.50	2
Yes	0.00	0.00	0.00	1
accuracy			0.33	3
macro avg	0.25	0.25	0.25	3
weighted avg	0.33	0.33	0.33	3

b. Build the model on train set and predict the values on test set

Code:-

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score

# Sample Data
df = pd.DataFrame({
    'tenure': [1, 2, 3, 4, 5, 6, 7, 8, 9, 10],
    'MonthlyCharges': [20, 30, 40, 50, 60, 70, 80, 90, 100, 110],
    'Churn': ['Yes', 'No', 'Yes', 'No', 'Yes', 'Yes', 'No', 'No', 'Yes', 'No']
})

# Prepare Data
X = df[['tenure', 'MonthlyCharges']]
y = df['Churn'].map({'Yes': 1, 'No': 0})
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, stratify=y, random_state=42)

# Train Model
clf = RandomForestClassifier(n_estimators=100, random_state=42)
clf.fit(X_train, y_train)

# Predict on Test Set
y_pred = clf.predict(X_test)

# Results
print(f"Predictions on Test Set: {y_pred}")
print(f"Accuracy: {accuracy_score(y_test, y_pred):.2f}")
```

Ans:-

```
Predictions on Test Set: [1 0 0]
Accuracy: 0.33
```

c. Build the confusion matrix and calculate the accuracy

Code:-

```
from sklearn.metrics import confusion_matrix, accuracy_score

# Example true labels and predicted labels
y_true = [1, 0, 1, 1, 0, 1, 0, 0, 1, 0] # Example true labels
y_pred = [1, 0, 1, 0, 0, 1, 0, 1, 1, 0] # Example predicted labels

# Generate the confusion matrix
conf_matrix = confusion_matrix(y_true, y_pred)
print("Confusion Matrix:")
print(conf_matrix)

# Calculate accuracy
accuracy = accuracy_score(y_true, y_pred)
print("\nAccuracy:", accuracy)
```

Confusion Matrix:
[[4 1]
 [1 4]]

Accuracy: 0.8

Ans:-

Confusion Matrix:- [[4 1]

[1 4]

Accuracy : 0.8